

First Full LoRA Trial with Transformer

peft (for LoRA) and FLAN-T5-small for the LLM

I'm following what seems to be a great tutorial from Mehul Gupta,

<https://medium.com/data-science-in-your-pocket/lora-for-fine-tuning-llms-explained-with-codes-and-example-62a7ac5a3578>

<https://web.archive.org/web/20240522140323/https://medium.com/data-science-in-your-pocket/lora-for-fine-tuning-llms-explained-with-codes-and-example-62a7ac5a3578>

I'm doing this to prepare creating a LoRA for RWKV (@todo put links in here) so as to fine-tune it for Pat's OLECT-LM stuff.

```
In [1]: # # Don't need this again
# !powershell -c (Get-Date -UFormat \"%s_%Y%m%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

1717091264_20240530T174744-0600

Output was:

1717091264_20240530T174744-0600

Imports

```
In [2]: from datasets import load_dataset
from random import randrange
import torch
from transformers import AutoTokenizer, \
                        AutoModelForSeq2SeqLM, \
                        TrainingArguments, \
                        pipeline
from transformers.utils import logging
from peft import LoraConfig, \
```

```

        prepare_model_for_kbit_training, \
        get_peft_model, \
        AutoPeftModelForCausalLM
from trl import SFTTrainer
from huggingface_hub import login, notebook_login

from datasets import load_metric
import nltk
import rouge_score

import pickle
import pprint
import timeit
from humanfriendly import format_timespan
import os

```

Load the training and test dataset along with the LLM with its tokenizer

The LLM will be fine-tuned. It seems the tokenizer will also be fine-tuned, but I'm not sure

Why aren't we loading the validation set? (I don't know; that's not a teaching question.)

I've tried to make use of it with the `trainer`. We'll see how it goes

```

In [3]: # Need to install datasets from pip, not conda. I'll do all from pip.
# I'll get rid of the current conda environment and make it anew.
# Actually, I'll make sure conda and pip are updated, then do what
# I discussed above.
#
# cf.
# arch_ref_1 = "https://web.archive.org/web/20240522150357/" + \
#             "https://stackoverflow.com/questions/77433096/" + \
#             "notimplementederror-loading-a-dataset-" + \
#             "cached-in-a-localfilesystem-is-not-suppor"
#
# Also useful might be
# arch_ref_2 = "https://web.archive.org/web/20240522150310/" + \
#             "https://stackoverflow.com/questions/76340743/" + \
#             "huggingface-load-datasets-gives-" + \
#             "notimplementederror-cannot-error"

```

```
#
data_files = {'train': 'samsum-train.json',
              'evaluation': 'samsum-validation.json',
              'test': 'samsum-test.json'}
dataset = load_dataset('json', data_files=data_files)

model_name = "google/flan-t5-small"
model = AutoModelForSeq2SeqLM.from_pretrained(model_name)

# Next line makes training faster but a little less accurate
model.config.pretraining_tp = 1

tokenizer = AutoTokenizer.from_pretrained(model_name,
                                         trust_remote_code=True)

# padding instructions for the tokenizer
#+ ??? !!! What about for RWKV !!! ???
#+ Will it be the same?
tokenizer.pad_token = tokenizer.eos_token
tokenizer.padding_side = "right"
```

```
Generating train split: 0 examples [00:00, ? examples/s]
Generating evaluation split: 0 examples [00:00, ? examples/s]
Generating test split: 0 examples [00:00, ? examples/s]
config.json: 0%|          | 0.00/1.40k [00:00<?, ?B/s]
```

C:\Users\Anast\.conda\envs\rwkv-lora-pat\lib\site-packages\huggingface_hub\file_download.py:157: UserWarning: `huggingface_hub` cache-system uses symlinks by default to efficiently store duplicated files but your machine does not support them in C:\Users\Anast\.cache\huggingface\hub\models--google--flan-t5-small. Caching files will still work but in a degraded version that might require more space on your disk. This warning can be disabled by setting the `HF_HUB_DISABLE_SYMLINKS_WARNING` environment variable. For more details, see https://huggingface.co/docs/huggingface_hub/how-to-cache#limitations.

To support symlinks on Windows, you either need to activate Developer Mode or to run Python as an administrator. In order to see activate developer mode, see this article: <https://docs.microsoft.com/en-us/windows/apps/get-started/enable-your-device-for-development>

```
warnings.warn(message)
```

```
model.safetensors: 0%|          | 0.00/308M [00:00<?, ?B/s]
generation_config.json: 0%|          | 0.00/147 [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/2.54k [00:00<?, ?B/s]
spiece.model: 0%|          | 0.00/792k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/2.42M [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/2.20k [00:00<?, ?B/s]
```

Trying some things I've been learning

```
In [4]: print(model)
```

```

T5ForConditionalGeneration(
  (shared): Embedding(32128, 512)
  (encoder): T5Stack(
    (embed_tokens): Embedding(32128, 512)
    (block): ModuleList(
      (0): T5Block(
        (layer): ModuleList(
          (0): T5LayerSelfAttention(
            (SelfAttention): T5Attention(
              (q): Linear(in_features=512, out_features=384, bias=False)
              (k): Linear(in_features=512, out_features=384, bias=False)
              (v): Linear(in_features=512, out_features=384, bias=False)
              (o): Linear(in_features=384, out_features=512, bias=False)
              (relative_attention_bias): Embedding(32, 6)
            )
            (layer_norm): T5LayerNorm()
            (dropout): Dropout(p=0.1, inplace=False)
          )
          (1): T5LayerFF(
            (DenseReluDense): T5DenseGatedActDense(
              (wi_0): Linear(in_features=512, out_features=1024, bias=False)
              (wi_1): Linear(in_features=512, out_features=1024, bias=False)
              (wo): Linear(in_features=1024, out_features=512, bias=False)
              (dropout): Dropout(p=0.1, inplace=False)
              (act): NewGELUActivation()
            )
            (layer_norm): T5LayerNorm()
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
    )
  )
  (1-7): 7 x T5Block(
    (layer): ModuleList(
      (0): T5LayerSelfAttention(
        (SelfAttention): T5Attention(
          (q): Linear(in_features=512, out_features=384, bias=False)
          (k): Linear(in_features=512, out_features=384, bias=False)
          (v): Linear(in_features=512, out_features=384, bias=False)
          (o): Linear(in_features=384, out_features=512, bias=False)
        )
        (layer_norm): T5LayerNorm()
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
)

```

```

    )
    (1): T5LayerFF(
      (DenseReluDense): T5DenseGatedActDense(
        (wi_0): Linear(in_features=512, out_features=1024, bias=False)
        (wi_1): Linear(in_features=512, out_features=1024, bias=False)
        (wo): Linear(in_features=1024, out_features=512, bias=False)
        (dropout): Dropout(p=0.1, inplace=False)
        (act): NewGELUActivation()
      )
      (layer_norm): T5LayerNorm()
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
)
(final_layer_norm): T5LayerNorm()
(dropout): Dropout(p=0.1, inplace=False)
)
(decoder): T5Stack(
  (embed_tokens): Embedding(32128, 512)
  (block): ModuleList(
    (0): T5Block(
      (layer): ModuleList(
        (0): T5LayerSelfAttention(
          (SelfAttention): T5Attention(
            (q): Linear(in_features=512, out_features=384, bias=False)
            (k): Linear(in_features=512, out_features=384, bias=False)
            (v): Linear(in_features=512, out_features=384, bias=False)
            (o): Linear(in_features=384, out_features=512, bias=False)
            (relative_attention_bias): Embedding(32, 6)
          )
          (layer_norm): T5LayerNorm()
          (dropout): Dropout(p=0.1, inplace=False)
        )
        (1): T5LayerCrossAttention(
          (EncDecAttention): T5Attention(
            (q): Linear(in_features=512, out_features=384, bias=False)
            (k): Linear(in_features=512, out_features=384, bias=False)
            (v): Linear(in_features=512, out_features=384, bias=False)
            (o): Linear(in_features=384, out_features=512, bias=False)
          )
          (layer_norm): T5LayerNorm()
        )
      )
    )
  )
)

```

```

        (dropout): Dropout(p=0.1, inplace=False)
    )
    (2): T5LayerFF(
      (DenseReluDense): T5DenseGatedActDense(
        (wi_0): Linear(in_features=512, out_features=1024, bias=False)
        (wi_1): Linear(in_features=512, out_features=1024, bias=False)
        (wo): Linear(in_features=1024, out_features=512, bias=False)
        (dropout): Dropout(p=0.1, inplace=False)
        (act): NewGELUActivation()
      )
      (layer_norm): T5LayerNorm()
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
)
(1-7): 7 x T5Block(
  (layer): ModuleList(
    (0): T5LayerSelfAttention(
      (SelfAttention): T5Attention(
        (q): Linear(in_features=512, out_features=384, bias=False)
        (k): Linear(in_features=512, out_features=384, bias=False)
        (v): Linear(in_features=512, out_features=384, bias=False)
        (o): Linear(in_features=384, out_features=512, bias=False)
      )
      (layer_norm): T5LayerNorm()
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (1): T5LayerCrossAttention(
      (EncDecAttention): T5Attention(
        (q): Linear(in_features=512, out_features=384, bias=False)
        (k): Linear(in_features=512, out_features=384, bias=False)
        (v): Linear(in_features=512, out_features=384, bias=False)
        (o): Linear(in_features=384, out_features=512, bias=False)
      )
      (layer_norm): T5LayerNorm()
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (2): T5LayerFF(
      (DenseReluDense): T5DenseGatedActDense(
        (wi_0): Linear(in_features=512, out_features=1024, bias=False)
        (wi_1): Linear(in_features=512, out_features=1024, bias=False)
        (wo): Linear(in_features=1024, out_features=512, bias=False)

```


RLHF ([archived wikipedia page](#)) is **R**einforcement **L**earning from **H**uman **F**eedback. TRL%20step.) ([archived](#)) **T**ransfer **R**einforcement **L**earning, a library from Hugging Face.

For the parameter, `formatting_func`, I can look at the documentation site above (specifically [here](#)), at the GitHub repo for [the code](#) (in the docstrings), or from my local `conda` environment, at `C:\Users\bballdave025\.conda\envs\rwkv-lora-pat\Lib\site-packages\trl\trainer\sft_trainer.py`.

Pulling code from the last one, I get

```
formatting_func (`Optional[Callable]`):
    The formatting function to be used for creating the `ConstantLengthDataset`.
```

That matches the first very well

```
formatting_func ( Optional[Callable] ) — The formatting function to be used for creating the
ConstantLengthDataset .
```

(A quick note: In this Jupyter Notebook environment, I could have typed `trainer = SFTTrainer(` and then `Shift` + `Tab` to find that same documentation.

However, I think that more clarity is found at the [documentation for `ConstantLengthDataset](#)

```
formatting_func ( Callable , optional ) — Function that formats the text before tokenization. Usually it is
recommended to have follows a certain pattern such as "### Question: {question} ### Answer: {answer}"
```

So, as we'll see the next code from the tutorial, it basically is a prompt templater/formatter that matches the JSON. For example, we use `sample['dialogue']` to access the `dialogue` key/pair. That's what I got from all this stuff.

Mehul Gupta himself stated

```
Next, using the Input and Output, we will create a prompt template which is a requirement by the SFTTrainer we will
be using later
```

Prompt

```
In [9]: def prompt_instruction_format(sample):
        return f""" Instruction:
            Use the Task below and the Input given to write the Response:

            ### Task:
            Summarize the Input

            ### Input:
            {sample['dialogue']}

            ### Response:
            {sample['summary']}
            """
        ##endof: prompt_instruction_format(sample)
```

Trainer - the LoRA Setup Part

Arguments and Configuration

```
In [12]: # Some arguments to pass to the trainer
training_args = TrainingArguments(
    output_dir='output',
    num_train_epochs=1,
    per_device_train_batch_size=4,
    save_strategy='epoch',
    learning_rate=2e-4,
    do_eval=True,
    per_device_eval_batch_size=4,
    eval_strategy='epoch',
    hub_model_id="dwb-flan-t5-small-lora-finetune",
)

# the fine-tuning (peft for LoRA) stuff
peft_config = LoraConfig( lora_alpha=16,
    lora_dropout=0.1,
    r=64,
    bias='none',
    task_type='CAUSAL_LM'
)
```

`task_type`, cf. <https://github.com/huggingface/peft/blob/main/src/peft/config.py#L222>

```

    Args:
        peft_type (Union[`~peft.utils.config.PeftType`], `str`): The type of Peft method to
        use.
        task_type (Union[`~peft.utils.config.TaskType`], `str`): The type of task to perform.
        inference_mode (`bool`, defaults to `False`): Whether to use the Peft model in
        inference mode.

```

After some searching using Cygwin

```

bballdave025@MYMACHINE /cygdrive/c/Users/bballdave025/.conda/envs/rwkv-lora-pat/Lib/site-
packages/peft/utils

```

```
$ ls -lah
```

```
total 116K
```

```

drwx-----+ 1 bballdave025 bballdave025    0 May 28 21:09 .
drwx-----+ 1 bballdave025 bballdave025    0 May 28 21:09 ..
-rwx-----+ 1 bballdave025 bballdave025 2.0K May 28 21:09 __init__.py
drwx-----+ 1 bballdave025 bballdave025    0 May 28 21:09 __pycache__
-rwx-----+ 1 bballdave025 bballdave025 8.0K May 28 21:09 constants.py
-rwx-----+ 1 bballdave025 bballdave025 3.8K May 28 21:09 integrations.py
-rwx-----+ 1 bballdave025 bballdave025 17K May 28 21:09 loftq_utils.py
-rwx-----+ 1 bballdave025 bballdave025 9.7K May 28 21:09 merge_utils.py
-rwx-----+ 1 bballdave025 bballdave025 25K May 28 21:09 other.py
-rwx-----+ 1 bballdave025 bballdave025 2.2K May 28 21:09 peft_types.py
-rwx-----+ 1 bballdave025 bballdave025 21K May 28 21:09 save_and_load.py

```

```

bballdave025@MYMACHINE /cygdrive/c/Users/bballdave025/.conda/envs/rwkv-lora-pat/Lib/site-
packages/peft/utils

```

```
$ grep -iIRHn "TaskType" .
```

```
peft_types.py:60:class TaskType(str, enum.Enum):
```

```
__init__.py:20:# from .config import PeftConfig, PeftType, PromptLearningConfig, TaskType
```

```
__init__.py:22:from .peft_types import PeftType, TaskType
```

```

bballdave025@MYMACHINE /cygdrive/c/Users/bballdave025/.conda/envs/rwkv-lora-pat/Lib/site-
packages/peft/utils

```

```
$
```

So, let's look at the `peft_types.py` file.

The docstring for `class TaskType(str, enum.Enum)` is

Enum class for the different types of tasks supported by PEFT.

Overview of the supported task types:

- SEQ_CLS: Text classification.
- SEQ_2_SEQ_LM: Sequence-to-sequence language modeling.
- CAUSAL_LM: Causal language modeling.
- TOKEN_CLS: Token classification.
- QUESTION_ANS: Question answering.
- FEATURE_EXTRACTION: Feature extraction. Provides the hidden states which can be used as embeddings or features for downstream tasks.

We're going to start timing stuff, so here's some system info

`win_system_info_as_script.py` is a script I wrote with the help of a variety of StackOverflow and documentation sources. It should be in the working directory.

```
In [13]: import win_system_info_as_script as winsysinfo
winsysinfo.run()
```

System Information

System: Windows
 Node Name: NOT-FOR-NOW
 Release: 10
 Version: 10.0.19045
 Machine: AMD64
 Processor: Intel64 Family 6 Model 165 Stepping 3, GenuineIntel
 Processor: Intel(R) Core(TM) i3-10100 CPU @ 3.60GHz
 Ip-Address: NOT-FOR-NOW
 Mac-Address: NOT-FOR-NOW

Boot Time

Boot Time (date and time of last boot) was
 Boot Time: 2024-5-26T14:29:0

CPU Info

Physical cores: 4
 Total cores: 8
 CPU Usage Per Core:
 Core 0: 3.1%
 Core 1: 1.6%
 Core 2: 6.2%
 Core 3: 0.0%
 Core 4: 3.1%
 Core 5: 0.0%
 Core 6: 6.2%
 Core 7: 0.0%
 Total CPU Usage: 5.7%
 Max Frequency: 3600.00Mhz
 Min Frequency: 0.00Mhz
 Current Frequency: 3600.00Mhz

GPU Info

Information on GPU(s)/Graphics Card(s)
 (if any such information is to be found)

Using wmi , we get the following win32_VideoController names.

Trigger 6 External Graphics

Intel(R) UHD Graphics 630

Using PyTorch and the torch.cuda.is_available() method.

The statement, 'There is CUDA and an appropriate GPU',

is ... False

Using TensorFlow with several of its methods.

Attempting to get GPU Device List

No GPU Devices.

Tensorflow can give us CPU (and/or GPU) info.

The info here might help you know if we're running on a CPU.

Trying to use some nvidia code (nvidia-smi) to find information

That's the end of the nvidia try.

Those are all our chances to find out about any GPU/Graphics Cards.

Memory (RAM) Information

Total: 31.67GbB

Available: 17.95GbB

Used: 13.72GbB

Percentage: 43.3%

===== SWAP Memory =====

That nvidia stuff didn't work

The error information is:

[WinError 2] The system cannot find the file specified

Total: 4.75GbB
Free: 4.64GbB
Used: 108.27MbB
Percentage: 2.2%

Disk Info

Partitions and Usage:

=== Device: C:\ ===

Mountpoint: C:\

File system type: NTFS

Total Size: 915.94GbB

Used: 587.01GbB

Free: 328.93GbB

Percentage: 64.1%

=== Device: D:\ ===

Mountpoint: D:\

File system type: exFAT

Total Size: 12.73TbB

Used: 1.99TbB

Free: 10.75TbB

Percentage: 15.6%

=== Device: E:\ ===

Mountpoint: E:\

File system type: FAT32

Total Size: 115.31GbB

Used: 46.08GbB

Free: 69.23GbB

Percentage: 40.0%

Since last boot,

Total read: 158.10GbB

Total write: 204.07GbB

ROUGE Metrics

Some references from the Microsoft/Google (who?) implementation

<https://pypi.org/project/rouge-score/>

<https://web.archive.org/web/20240530231357/https://pypi.org/project/rouge-score/>

<https://github.com/google-research/google-research/tree/master/rouge>

<https://web.archive.org/web/20240530231412/https://github.com/google-research/google-research/tree/master/rouge>

Not the one I used:

https://github.com/microsoft/nlp-recipes/blob/master/examples/text_summarization/summarization_evaluation.ipynb

https://web.archive.org/web/20240530231709/https://github.com/microsoft/nlp-recipes/blob/master/examples/text_summarization/summarization_evaluation.ipynb

Someone else made this other one, which I inspected but didn't use.

<https://pypi.org/project/rouge/>

<https://web.archive.org/web/20240530232029/https://pypi.org/project/rouge/>

<https://github.com/pltrdy/rouge>

<https://web.archive.org/web/20240530232023/https://github.com/pltrdy/rouge>

but I think he defers to the rouge_score from Google.

My ROUGE Metrics

I want to use the skip-grams score. Thanks to

<https://www.bomberbot.com/machine-learning/skip-bigrams-in-system/>

<https://web.archive.org/web/20240530230949/https://www.bomberbot.com/machine-learning/skip-bigrams-in-system/>

I can do this as well as writing the code for the other metrics.


```

In [14]: # import itertools

# def rouge_n(system, reference, n):
#     '''
#     # ROUGE-N : N-Grams implementation

#     # ref = "https://web.archive.org/web/20240530230949/" + \
#         # "https://www.bomberbot.com/machine-learning/" + \
#         # "skip-bigrams-in-system/"

#     @param system      string    The hypothesis
#     @param reference    string    The truth
#     @param n            string    The "n" in "n-gram", i.e.
#                                   # the number of words in
#                                   # each grouping

#     @returns dict in form {"recall": recall,
#                             # "precision": precision,
#                             # "f-measure": f_measure}

#     Example:
#     # >>> import rouge_n
#     # >>>
#     # >>> system = "The cat was found under the bed."
#     # >>> reference = "The cat was hidden under the bed."
#     # >>>
#     # >>> print(rouge_n(system, reference, 1)) # ROUGE-1
#     # >>> print(rouge_n(system, reference, 2)) # ROUGE-2
#     # {'recall': 0.8571428571428571, 'precision': 1.0, 'f-measure': 0.9230769230769231}
#     # {'recall': 0.6, 'precision': 0.5, 'f-measure': 0.5454545454545455}
#     '''

#     sys_ngrams = list(itertools.ngrams(system.split(), n))
#     ref_ngrams = list(itertools.ngrams(reference.split(), n))

#     overlaps = set(sys_ngrams) & set(ref_ngrams)
#     recall = len(overlaps) / len(ref_ngrams)
#     precision = len(overlaps) / len(sys_ngrams)

#     if precision + recall == 0:
#         f_measure = 0

```

```

# else:
#     # f_measure = 2 * precision * recall / (precision + recall)
# ##endof: if/else precision + recall == 0
# return {"recall": recall, "precision": precision, "f-measure": f_measure}
# ##endof: rouge_n(system, reference, n)

# def lcs(X, Y):
#     '''
#     # Longest common subsequence
#     # '''

#     # m = len(X)
#     # n = len(Y)

#     # L = [[None]*(n+1) for i in range(m+1)]

#     # for i in range(m+1):
#     #     # for j in range(n+1):
#     #         # if i == 0 or j == 0:
#     #             # L[i][j] = 0
#     #         # elif X[i-1] == Y[j-1]:
#     #             # L[i][j] = L[i-1][j-1]+1
#     #         # else:
#     #             # L[i][j] = max(L[i-1][j], L[i][j-1])
#     #     ##endof: if
#     # ##endof: for j
#     # ##endof: for i
#     # return L[m][n]
# ##endof: lcs(X, Y)

# def rouge_l(system, reference):
#     '''
#     # ROUGE-L : Longest Common Subsequence implementation

#     # ref = "https://web.archive.org/web/20240530230949/" + \
#     # "https://www.bomberbot.com/machine-learning/" + \
#     # "skip-bigrams-in-system/"

#     # @param system      string      The hypothesis
#     # @param reference    string      The truth

```

```

# @returns dict in form {"recall": recall,
                        # "precision": precision,
                        # "f-measure": f_measure}

# Example:
# >>> import rouge_l, lcs
# >>>
# >>> system = "The quick dog jumps over the lazy fox."
# >>> reference = "The quick brown fox jumps over the lazy dog."
# >>>
# >>> print(rouge_l(system, reference))
# {'recall': 0.7777777777777778, 'precision': 0.875, 'f-measure': 0.823529411764706}
# '''

# sys_len = len(system.split())
# ref_len = len(reference.split())
# lcs_len = lcs(system.split(), reference.split())

# recall = lcs_len / ref_len
# precision = lcs_len / sys_len

# if precision + recall == 0:
#     # f_measure = 0
# else:
#     # f_measure = 2 * precision * recall / (precision + recall)
# ##endof: if/else

# return {"recall": recall, "precision": precision, "f-measure": f_measure}
# ##endof: rouge_l(system, reference)

# from itertools import combinations

# def skipbigrams(sequence, n):
#     '''
#     # Returns the set of skip n-grams

#     # @param sequence
#     # @param n

```

```

# '''

# return set(combinations(sequence, n))

# ##endof: skipbigrams(sequence, n=2)

# def rouge_s(system, reference, n=2):
# '''
# ROUGE-S : Skip Bigrams implementation

# @param
# @param
# @param

# @returns

# Example
# >>> import skipbigrams, rouge_s
# >>>
# >>> system = "The quick dog jumps over the lazy fox."
# >>> reference = "The quick brown fox jumps over the lazy dog."
# >>>
# >>> print(rouge_s(system, reference))
# {'recall': 0.35, 'precision': 0.4166666666666667, 'f-measure': 0.38095238095238093}
# '''
# sys_skipbigrams = skipbigrams(system.split(), n)
# ref_skipbigrams = skipbigrams(reference.split(), n)

# overlaps = sys_skipbigrams & ref_skipbigrams
# recall = len(overlaps) / len(ref_skipbigrams)
# precision = len(overlaps) / len(sys_skipbigrams)

# if precision + recall == 0:
#     f_measure = 0
# else:
#     f_measure = 2 * precision * recall / (precision + recall)
# ##endof: if/else

# return {"recall": recall, "precision": precision, "f-measure": f_measure}
# ##endof: rouge_s(system, reference, n=2)

```

Try for a baseline

Just one summarization to begin with, randomly picked

```
In [15]: # # Don't need this again
#!powershell -c (Get-Date -UFormat \"%s_%Y%m%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

1717094554_20240530T184234-0600

Output was:

```
1717094554_20240530T184234-0600
```

```
In [17]: # Just one summarization to begin with, randomly picked ... but
# now with th possibility of a known seed, to allow visual
# comparison with after-training results.
# I'M NOT GOING TO USE THIS REPEATED SEED, I'm just going to
# use the datum at the first index to compare.
```

```
do_seed_for_repeatable = False

summarizer = pipeline('summarization',
                      model=model,
                      tokenizer=tokenizer)

if do_seed_for_repeatable:
    rand_seed_for_randrange = 137
    random.seed(rand_seed_for_randrange)
##endof: if do_seed_for_repeatable

sample = dataset['test'][randrange(len(dataset["test"]))]
print(f"dialogue: \n{sample['dialogue']}\n-----")

res = summarizer(sample["dialogue"])

print(f"flan-t5-small summary:\n{res[0]['summary_text']}")
```

Your max_length is set to 200, but your input_length is only 122. Since this is a summarization task, where outputs shorter than the input are typically wanted, you might consider decreasing max_length manually, e.g. summarizer('...', max_length=61)

dialogue:

Harry: and? have you listened to it?

Jacob: listened to what?

Harry: to the song i sent you 3 days ago --

Jacob: oh shit, i completely forgot...

Harry: ofc again

Jacob: don't be like this :* i'll do that later tonight

Harry: heh, okay

Harry: i'm really curious what you'll think about it

Jacob: i'll let you know, a bit busy right now, speak to you later!

Harry: okay

flan-t5-small summary:

Jacob forgot to listen to the song he sent Jacob 3 days ago. Harry will let Jacob know later tonight. Jacob will talk to Harry later.

Now, one summarization with comparison to ground truth

```
In [18]: summarizer = pipeline('summarization',
                                model=model,
                                tokenizer=tokenizer)

pred_test_list = []
ref_test_list = []

sample_num = 0

this_sample = dataset['test'][sample_num]

print(f"dialogue: \n{this_sample['dialogue']}\n-----")

grnd_summary = this_sample['summary']
res = summarizer(this_sample['dialogue'])
res_summary = res[0]['summary_text']

# humgen is for human-generated

print(f"human-genratd summary:\n{grnd_summary}")
print(f"flan-t5-small summary:\n{res_summary}")

ref_test_list.append(grnd_summary)
```

```

pred_test_list.append(res_summary)

print("\n\n----- ROUGE SCORES -----")

rouge = load_metric('rouge', trust_remote_code=True)

results = rouge.compute(predictions=pred_test_list,
                        references=ref_test_list,
                        use_aggregator=True)

# >>> print(list(results.keys()))
# ['rouge1', 'rouge2', 'rougeL', 'rougeLsum']

print()
print("ROUGE-1 results")
pprint.pp(results['rouge1'])
print()
print("ROUGE-2 results")
pprint.pp(results['rouge2'])
print()
print("ROUGE-L results")
pprint.pp(results['rougeL'])
print()
print("ROUGE-Lsum results")
pprint.pp(results['rougeLsum'])

```

Your max_length is set to 200, but your input_length is only 133. Since this is a summarization task, where outputs shorter than the input are typically wanted, you might consider decreasing max_length manually, e.g. summarizer('...', max_length=66)

dialogue:

Hannah: Hey, do you have Betty's number?

Amanda: Lemme check

Hannah: <file_gif>

Amanda: Sorry, can't find it.

Amanda: Ask Larry

Amanda: He called her last time we were at the park together

Hannah: I don't know him well

Hannah: <file_gif>

Amanda: Don't be shy, he's very nice

Hannah: If you say so..

Hannah: I'd rather you texted him

Amanda: Just text him 😊

Hannah: Urgh.. Alright

Hannah: Bye

Amanda: Bye bye

human-genratd summary:

Hannah needs Betty's number but Amanda doesn't have it. She needs to contact Larry.

flan-t5-small summary:

Larry called Hannah last time she was at the park together. Hannah doesn't know Larry well. Larry called her last time they were at a park. Hannah will text Larry.

----- ROUGE SCORES -----

C:\Users\Anast\AppData\Local\Temp\ipykernel_8400\2790818686.py:26: FutureWarning: load_metric is deprecated and will be removed in the next major version of datasets. Use 'evaluate.load' instead, from the new library 😊 Evaluate: <https://huggingface.co/docs/evaluate>

rouge = load_metric('rouge', trust_remote_code=True)

Downloading builder script: 0%| | 0.00/2.17k [00:00<?, ?B/s]

ROUGE-1 results

```
AggregateScore(low=Score(precision=0.16129032258064516, recall=0.3125, fmeasure=0.2127659574468085), mid=Score(precision=0.16129032258064516, recall=0.3125, fmeasure=0.2127659574468085), high=Score(precision=0.16129032258064516, recall=0.3125, fmeasure=0.2127659574468085))
```

ROUGE-2 results

```
AggregateScore(low=Score(precision=0.03333333333333333, recall=0.06666666666666667, fmeasure=0.04444444444444444), mid=Score(precision=0.03333333333333333, recall=0.06666666666666667, fmeasure=0.04444444444444444), high=Score(precision=0.03333333333333333, recall=0.06666666666666667, fmeasure=0.04444444444444444))
```

ROUGE-L results

```
AggregateScore(low=Score(precision=0.12903225806451613, recall=0.25, fmeasure=0.1702127659574468), mid=Score(precision=0.12903225806451613, recall=0.25, fmeasure=0.1702127659574468), high=Score(precision=0.12903225806451613, recall=0.25, fmeasure=0.1702127659574468))
```

ROUGE-Lsum results

```
AggregateScore(low=Score(precision=0.12903225806451613, recall=0.25, fmeasure=0.1702127659574468), mid=Score(precision=0.12903225806451613, recall=0.25, fmeasure=0.1702127659574468), high=Score(precision=0.12903225806451613, recall=0.25, fmeasure=0.1702127659574468))
```

Verbosity stuff - get rid of the nice advice

```
In [19]: # # Don't need this again
# !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

```
1717094688_2024-05-30T184448-0600
```

Output was:

```
1717094688_2024-05-30T184448-0600
```

```
In [20]: log_verbosity_is_critical = \
    logging.get_verbosity() == logging.CRITICAL # alias FATAL, 50
log_verbosity_is_error = \
    logging.get_verbosity() == logging.ERROR # 40
log_verbosity_is_warn = \
    logging.get_verbosity() == logging.WARNING # alias WARN, 30
log_verbosity_is_info = \
    logging.get_verbosity() == logging.INFO # 20
log_verbosity_is_debug = \
    logging.get_verbosity() == logging.DEBUG # 10
```

```

print( "The statement, 'logging verbosity is CRITICAL' " + \
      f"is {log_verbosity_is_critical}")
print( "The statement, 'logging verbosity is    ERROR' " + \
      f"is {log_verbosity_is_error}")
print( "The statement, 'logging verbosity is  WARNING' " + \
      f"is {log_verbosity_is_warn}")
print( "The statement, 'logging verbosity is    INFO' " + \
      f"is {log_verbosity_is_info}")
print( "The statement, 'logging verbosity is  DEBUG' " + \
      f"is {log_verbosity_is_debug}")

print()

init_log_verbosity = logging.get_verbosity()
print(f"The value of logging.get_verbosity() is: {init_log_verbosity}")

print()

init_t_n_a_w = os.environ.get('TRANSFORMERS_NO_ADVISORY_WARNINGS')
print(f"TRANSFORMERS_NO_ADIVSORY_WARNINGS: {init_t_n_a_w}")

```

The statement, 'logging verbosity is CRITICAL' is False
 The statement, 'logging verbosity is ERROR' is False
 The statement, 'logging verbosity is WARNING' is True
 The statement, 'logging verbosity is INFO' is False
 The statement, 'logging verbosity is DEBUG' is False

The value of logging.get_verbosity() is: 30

TRANSFORMERS_NO_ADIVSORY_WARNINGS: None

Actual Baseline

```

In [21]: # # Don't need this again
        # !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'

```

1717094729_2024-05-30T184529-0600

Output was:

1717094729_2024-05-30T184529-0600

```

In [22]: # ref1 = "https://web.archive.org/web/20240530051418/" + \
#+         "https://stackoverflow.com/questions/73221277/" + \
#+         "python-hugging-face-warning"
# ref2 = "https://web.archive.org/web/20240530051559/" + \
#+       "https://huggingface.co/docs/transformers/en/" + \
#+       "main_classes/logging"

## Haven't tried this, because the logging seemed easier,
##+ and the logging worked
#os.environ("TRANSFORMERS_NO_ADVISORY_WARNINGS") = 1

logging.set_verbosity_error()

summarizer = pipeline('summarization',
                      model=model,
                      tokenizer=tokenizer)

prediction_list = []
reference_list = []

baseline_tic = timeit.default_timer()

for sample_num in range(len(dataset['test'])):
    this_sample = dataset['test'][sample_num]

    #print(f"dialogue: \n{this_sample['dialogue']}\n-----")

    grnd_summary = this_sample['summary']
    res = summarizer(this_sample['dialogue'])
    res_summary = res[0]['summary_text']

    #print(f"human-genratd summary:\n{grnd_summary}")
    #print(f"flan-t5-small summary:\n{res_summary}")

    reference_list.append(grnd_summary)
    prediction_list.append(res_summary)
##endof: for sample_num in range(len(dataset['test']))

baseline_toc = timeit.default_timer()

baseline_duration = baseline_toc - baseline_tic

```

```

print( "Getting things ready for scoring")
print(f"took {baseline_toc - baseline_tic:0.4f} seconds.")

print("\n\n----- ROUGE SCORES -----")

# @todo : Load it straight from the python package

rouge = load_metric('rouge', trust_remote_code=True)
# Set trust_remote_code=False to see the warning,
#+ deprecation, and what to change to.

results = rouge.compute(predictions=prediction_list,
                        references=reference_list,
                        use_aggregator=True)

# >>> print(List(results.keys()))
# ['rouge1', 'rouge2', 'rougeL', 'rougeLsum']

print()
print("ROUGE-1 results")
pprint.pp(results['rouge1'])
print()
print("ROUGE-2 results")
pprint.pp(results['rouge2'])
print()
print("ROUGE-L results")
pprint.pp(results['rougeL'])
print()
print("ROUGE-Lsum results")
pprint.pp(results['rougeLsum'])

## Haven't tried this, because the logging seemed easier,
##+ and the logging worked
# os.environ("TRANSFORMERS_NO_ADVISORY_WARNINGS") = init_t_n_a_w

logging.set_verbosity(init_log_verbosity)

```

Getting things ready for scoring
took 1162.5236 seconds.

----- ROUGE SCORES -----

ROUGE-1 results

```
AggregateScore(low=Score(precision=0.3623620420489957, recall=0.5387757354848512, fmeasure=0.4120367260545498), mid=Score(precision=0.37354505712494124, recall=0.5519205917207157, fmeasure=0.42157981166777414), high=Score(precision=0.38488859010129967, recall=0.5656367969330082, fmeasure=0.4313347594247768))
```

ROUGE-2 results

```
AggregateScore(low=Score(precision=0.15911356172159186, recall=0.2428538196441747, fmeasure=0.18143374370228235), mid=Score(precision=0.16776558103745187, recall=0.256702553043478, fmeasure=0.1902802753555807), high=Score(precision=0.176789440182549, recall=0.26994546288441695, fmeasure=0.19927249226979166))
```

ROUGE-L results

```
AggregateScore(low=Score(precision=0.2804766356825748, recall=0.4221646517546973, fmeasure=0.31994624442237346), mid=Score(precision=0.2892456873611609, recall=0.43475844856226864, fmeasure=0.32792714750993834), high=Score(precision=0.2984044560730355, recall=0.44750618279601273, fmeasure=0.33706771629160553))
```

ROUGE-Lsum results

```
AggregateScore(low=Score(precision=0.28022879848567583, recall=0.4220608213922142, fmeasure=0.319655889478565), mid=Score(precision=0.289591524472137, recall=0.43469242939551495, fmeasure=0.32826109350219757), high=Score(precision=0.29852002937109373, recall=0.44804213486471733, fmeasure=0.33687287121399573))
```

```
In [24]: do_enter_duration_manually = False
NUM_TO_CATCH_NO_MANUAL_ENTRY = -137.
is_a_manual_entry_skip = False # innocent until proven guilty

if do_enter_duration_manually:
    # !!! remember to type in your number, if needed !!! #
    baseline_duration = NUM_TO_CATCH_NO_MANUAL_ENTRY
    # !!! UNCOMMENT THE NEXT LINE IF YOU WANT TO ENTER MANUALLY !!!
    #baseline_duration =
##endof: if do_enter_duration_manually

print("Running baseline inference (using the test set)")
if ( ( do_enter_duration_manually ) and \
    ( baseline_duration == -137. ) \
):
```

```

    print("took AN UNKNOWN AMOUNT OF TIME.")
    print("You didn't manually enter in your real time,")
    print("as you should have.")
    is_a_manual_entry_skip = True
elif ( ( do_enter_duration_manually ) and \
      ( baseline_duration != -137. )
      ):
    print("(and using your manually entered time)")
else:
    pass
##endof:  if <check manual entry

if not is_a_manual_entry_skip:
    print(f"took {format_timespan(baseline_duration)}")
##endof:  if not is_a_manual_entry_skip

```

Running baseline inference (using the test set)
 took 19 minutes and 22.52 seconds

Trainer - the Actual Trainer Part

```

In [25]: # # Don't need this again
        # !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'

```

1717096214_2024-05-30T191014-0600

Output was:

1717096214_2024-05-30T191014-0600

```

In [27]: trainer = SFTTrainer( model=model,
                             train_dataset=dataset['train'],
                             eval_dataset=dataset['evaluation'],
                             peft_config=peft_config,
                             tokenizer=tokenizer,
                             packing=True,
                             formatting_func=prompt_instruction_format,
                             args=training_args,
                             )
        ## Warnings are below output.

```

```
## Ended up not using this.
#                               max_seq_length=675
#                               )
```

WARNING:bitsandbytes.cextension:The installed version of bitsandbytes was compiled without GPU support. 8-bit optimizers, 8-bit multiplication, and GPU quantization are unavailable.

C:\Users\Anast\.conda\envs\rwkv-lora-pat\lib\site-packages\trl\trainer\sft_trainer.py:246: UserWarning: You didn't pass a `max_seq_length` argument to the SFTTrainer, this will default to 512

```
warnings.warn(
```

Generating train split: 0 examples [00:00, ? examples/s]

Generating train split: 0 examples [00:00, ? examples/s]

First time warnings from the code above (as it still is).

```
WARNING:bitsandbytes.cextension:The installed version of bitsandbytes \
was compiled without GPU support. 8-bit optimizers, 8-bit multiplication, \
and GPU quantization are unavailable.
C:\Users\bballldave025\.conda\envs\rwkv-lora-pat\lib\site-packages\trl\
trainer\sft_trainer.py:246: UserWarning: You didn't pass a `max_seq_length` \
argument to the SFTTrainer, this will default to 512
warnings.warn(
```

```
[ > Generating train split: 6143/0 [00:04<00:00, 2034.36 examples/s] ]
```

```
Token indices sequence length is longer than the specified maximum sequence \
length for this model (657 > 512). Running this sequence through the model \
will result in indexing errors
```

```
[ > Generating train split: 355/0 [00:00<00:00, 6.10 examples/s] ]
```

DWB Note

So, I'm changing the `max_seq_length` : Maybe I should just throw out the offender(s) (along with the blank one that's in there somewhere), but I'll just continue as is.

Actually, it appears I didn't run the updated cell, (with `max_seq_length=675`), since the Warning and Advice are still there.

Let's Train This LoRA Thing and See How It Does!

```
In [28]: # # Don't need this again
# !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

1717096271_2024-05-30T191111-0600

Output was:

1717096271_2024-05-30T191111-0600

At about 1717063394_2024-05-30T100314-0600, DWB went in and renamed profile.ps1 to NOT-USING_-_pro_file_-_now.ps1.bak That should get rid of our errors from powershell

The long-time-taking training code is just below.

```
In [29]: tic = timeit.default_timer()
trainer.train()
toc = timeit.default_timer()
print(f"tic: {tic}")
print(f"toc: {toc}")
training_duration = toc - tic
print(f"Training took {toc - tic:0.4f} seconds.")
```

 [1536/1536 3:04:33, Epoch 1/1]

Epoch	Training Loss	Validation Loss
-------	---------------	-----------------

1	0.068500	0.022573
---	----------	----------

C:\Users\Anast\.conda\envs\rwkv-lora-pat\lib\site-packages\huggingface_hub\file_download.py:1132: FutureWarning: `resume_download` is deprecated and will be removed in version 1.0.0. Downloads always resume when possible. If you want to force a new download, use `force_download=True`.

```
warnings.warn(
tic: 362634.7966071
toc: 373716.499057
Training took 11081.7024 seconds.
```

```
In [30]: do_by_hand = False
NUM_TO_CATCH_NO_DO_BY_HAND = -137.
is_a_do_by_hand_skip = False # innocent until proven guilty
```



```

if do_by_hand:
    # !!! remember to type in your number, if needed !!! #
    training_duration = NUM_TO_CATCH_NO_MANUAL_ENTRY
    # !!! UNCOMMENT THE NEXT LINE IF YOU WANT TO ENTER MANUALLY !!!
    #training_duration = 11081.7024
##endof:  if do_by_hand

print("Running training (using the training and eval sets)")
if ( ( do_by_hand ) and \
    ( training_duration == -137. ) \
):
    print("took AN UNKNOWN AMOUNT OF TIME.")
    print("You didn't manually enter in your real time,")
    print("as you should have.")
    is_a_do_by_hand_skip = True
elif ( ( do_by_hand ) and \
    ( training_duration != -137. )
):
    print("(and using your manually entered time)")
else:
    pass
##endof:  if <check manual entry

if not is_a_do_by_hand_skip:
    print(f"took {format_timespan(training_duration)}")
##endof:  if not is_a_manual_entry_skip

```

Training with LoRA (and with the other info as above)
 took 3 hours, 4 minutes and 41.7 seconds.

```

In [31]: # # Don't need this again
        # !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'

```

1717107458_2024-05-30T221738-0600

Output was:

```
1717107458_2024-05-30T221738-0600
```

@todo : consolidate "the other info as above"

I'm talking about the numbers of data points, tokens, whatever.

Any Comments / Things to Try (?)

We passed an evaluation set (parameter ``) to the `trainer`. How can we see information about that?

How to get the evaluation set used by the trainer

I added the following parameters to the `training_args = TrainingArguments(<args>)` call.

- `do_eval=True`
- `per_device_eval_batch_size=4`
- `eval_strategy='epoch'`

How to specify your repo name

I also added this next parameter to the arguments for `training_args = TrainingArguments(<args>)`

- `hub_model_id="dwb-flan-t5-small-lora-finetune"`

The final TrainingArguments call - with parameter list

```
training_args = TrainingArguments(  
    output_dir='output',  
    num_train_epochs=1,  
    per_device_train_batch_size=4,  
    save_strategy='epoch',  
    learning_rate=2e-4,  
    do_eval=True,  
    per_device_eval_batch_size=4,  
    eval_strategy='epoch',  
    hub_model_id="dwb-flan-t5-small-lora-finetune",  
)
```

Save the Trainer to Hugging Face and Get Our Updated Model

```
In [34]: # # Don't need this again
# !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

1717145367_2024-05-31T084927-0600

Output was:

```
1717145367_2024-05-31T084927-0600
```

I'm following the [\(archived\) tutorial from Mehul Gupta on Medium](#); since it's archived, you can follow exactly what I'm doing.

Running this next line of code will come up with a dialog box with text entry, and I'm now using the `@thebballdave025` for Hugging Face stuff.

Make sure to use the WRITE token, here.

```
In [35]: # This will come up with a dialog box with text entry,
# and I'm now using
# ( @thebballdave025 for Hugging Face ) HF stuff.

# Use the write token, here.
notebook_login()
```

VBox(children=(HTML(value='<center> <img\\nsrc=https://huggingface.co/front/assets/huggingface_logo-noborder.sv...

```
In [36]: # Save tokenizer and create a tokenizer model card
tokenizer.save_pretrained('testing')
# 'testing' is the local directory

# Create the trainer model card
trainer.create_model_card()

# Push the results to the Hugging Face Hub
trainer.push_to_hub()
```

C:\Users\Anast\.conda\envs\rwkv-lora-pat\lib\site-packages\huggingface_hub\file_download.py:1132: FutureWarning: `resume_download` is deprecated and will be removed in version 1.0.0. Downloads always resume when possible. If you want to force a new download, use `force_download=True`.

```
warnings.warn(
adapter_model.safetensors: 0%|          | 0.00/11.0M [00:00<?, ?B/s]
events.out.tfevents.1717084743.DESKTOP-07KM5A5.8272.0: 0%|          | 0.00/12.3k [00:00<?, ?B/s]
```

```

training_args.bin:  0%|          | 0.00/5.11k [00:00<?, ?B/s]
events.out.tfevents.1717117975.DESKTOP-07KM5A5.8400.0:  0%|          | 0.00/7.14k [00:00<?, ?B/s]
Upload 4 LFS files:  0%|          | 0/4 [00:00<?, ?it/s]

```

```

Out[36]: CommitInfo(commit_url='https://huggingface.co/thebballdave025/dwb-flan-t5-small-lora-finetune/commit/c87d34b398f3801
ceb1e18c819a7c8fc894989c7', commit_message='End of training', commit_description='', oid='c87d34b398f3801ceb1e18c819
a7c8fc894989c7', pr_url=None, pr_revision=None, pr_num=None)

```

Part of the output included the URL,

<https://huggingface.co/thebballdave025/dwb-flan-t5-small-lora-finetune/commit/c87d34b398f3801ceb1e18c819a7c8fc894989c7>

Hooray! The repo name I used in constructing the trainer worked!

I can get to the general repo with the URL,

<https://huggingface.co/thebballdave025/dwb-flan-t5-small-lora-finetune>

Info on the Fine-Tuned Model from the Repo's README - Model Card(?)

[thebballdave025/dwb-flan-t5-small-lora-finetune](#)

~~archived~~ The archiving attempt at archive.org (Wayback Machine) failed. I'm not sure why, as the model is set as public.

PEFT	TensorBoard	Safetensors	generator	trl	sft	generated_from_trainer	License: apache-2.0
------	-------------	-------------	-----------	-----	-----	------------------------	---------------------

@todo : [Edit Model Card](#)

Unable to determine this model's pipeline type. Check the docs [\(i\)](#).

Adapter for [google/flan-t5-small](#)

dwb-flan-t5-small-lora-finetune

This model is a fine-tuned version of [google/flan-t5-small](#) on the generator dataset [DWB note: I don't know why it says "generator dataset". I used the Samsum dataset, which I will link here and on the model card, eventually].

It achieves the following results on the evaluation set:

- Loss: 0.0226
- *DWB Note: I don't know which metric was used to calculate loss. If this were more important, I'd dig through code to find out and evaluate with the same metric. If I'm really lucky, they somehow used the ROUGE scores in the loss function, so we match.*

Model description

More information needed

Intended uses & limitations

More information needed

Training and evaluation data

More information needed

Training procedure

Training hyperparameters

The following hyperparameters were used during training:

- learning_rate: 0.0002
- train_batch_size: 4
- eval_batch_size: 4
- seed: 42
- optimizer: Adam with betas=(0.9,0.999) and epsilon=1e-08
- lr_scheduler_type: linear
- num_epochs: 1

Training results

Training Loss	Epoch	Step	Validation Loss
0.0685	1.0	1536	0.0226

Framework versions

- PEFT 0.11.2.dev0
- Transformers 4.41.1
- Pytorch 2.3.0+cpu
- Datasets 2.19.1
- Tokenizers 0.19.1

Actually Get the Model from Hugging Face

Running this next line of code will come up with a dialog box with text entry, and I'm now using the `@thebballdave025` for Hugging Face stuff.

Make sure to use the READ token, here.

```
In [ ]: ## Not run yet for this save.

# Read token. Will bring up text entry to paste token string
#notebook_login()
```

```
In [ ]:
```

```
In [ ]:
```

Evaluation on the Test Set and Comparison to Baseline

Verbosity stuff - get rid of the nice advice

```
In [ ]: !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

Output was:

```
timestamp
```

```
In [ ]: # bballdave025@MYMACHINE /cygdrive/c/Users/bballdave025/.conda/envs/rwkv-lora-pat/Lib/site-packages/peft/utis
# $ date +%s_%Y-%m-%dT%H%M%S%z'
# 1717049876_2024-05-30T001756-0600

log_verbosity_is_critical = \
    logging.get_verbosity() == logging.CRITICAL # alias FATAL, 50
log_verbosity_is_error = \
    logging.get_verbosity() == logging.ERROR # 40
log_verbosity_is_warn = \
    logging.get_verbosity() == logging.WARNING # alias WARN, 30
log_verbosity_is_info = \
    logging.get_verbosity() == logging.INFO # 20
log_verbosity_is_debug = \
    logging.get_verbosity() == logging.DEBUG # 10

print( "The statement, 'logging verbosity is CRITICAL' " + \
    f"is {log_verbosity_is_critical}")
print( "The statement, 'logging verbosity is ERROR' " + \
    f"is {log_verbosity_is_error}")
print( "The statement, 'logging verbosity is WARNING' " + \
    f"is {log_verbosity_is_warn}")
print( "The statement, 'logging verbosity is INFO' " + \
    f"is {log_verbosity_is_info}")
print( "The statement, 'logging verbosity is DEBUG' " + \
    f"is {log_verbosity_is_debug}")

print()

init_log_verbosity = logging.get_verbosity()
print(f"The value of logging.get_verbosity() is: {init_log_verbosity}")

print()
```

```
init_t_n_a_w = os.environ.get('TRANSFORMERS_NO_ADVISORY_WARNINGS')
print(f"TRANSFORMERS_NO_ADIVSORY_WARNINGS: {init_t_n_a_w}")
```

Here's the actual evaluation

```
In [ ]: !powershell -c (Get-Date -UFormat \"%s_%Y-%m-%dT%H%M%S%Z00\") -replace '[.][0-9]*_', '_'
```

Output was:

```
timestamp
```

!!! NOTE !!! I'm going to use `tat` (with an underscore or undescores before, after, or surrounding the variable names) to indicate 'testing-after-training'.

```
In [ ]: # I'm going to use 'tat' for testing-after-training

logging.set_verbosity_error()

summarizer = pipeline('summarization',
                      model=tat_model,
                      tokenizer=tokenizer)

prediction_tat_list = []
reference_tat_list = []

tat_tic = timeit.default_timer()

for sample_num in range(len(dataset['test'])):
    this_sample = dataset['test'][sample_num]

    #print(f"dialogue: \n{this_sample['dialogue']}\n-----")

    grnd_tat_summary = this_sample['summary']
    res_tat = summarizer(this_sample['dialogue'])
    res_tat_summary = res_tat[0]['summary_text']

    #print(f"human-genratd summary:\n{grnd_tat_summary}")
    #print(f"flan-t5-small summary:\n{res_tat_summary}")
```



```

reference_tat_list.append(grnd_tat_summary)
prediction_tat_list.append(res_tat_summary)
##endof: for sample_num in range(len(dataset['test']))

tat_toc = timeit.default_timer()

print( "Getting things ready for scoring (after training)")
print(f"took {tat_toc - tat_tic:0.4f} seconds.")

print("\n\n----- ROUGE SCORES -----")

rouge = load_metric('rouge', trust_remote_code=True)
# Set trust_remote_code=False to see the warning,
#+ deprecation, and what to change to.

results_tat = rouge.compute(
    predictions=prediction_tat_list,
    references=reference_tat_list,
    use_aggregator=True
)

# >>> print(List(results_tat.keys()))
# ['rouge1', 'rouge2', 'rougeL', 'rougeLsum']

print()
print("ROUGE-1 results")
pprint.pp(results_tat['rouge1'])
print()
print("ROUGE-2 results")
pprint.pp(results_tat['rouge2'])
print()
print("ROUGE-L results")
pprint.pp(results_tat['rougeL'])
print()
print("ROUGE-Lsum results")
pprint.pp(results_tat['rougeLsum'])

logging.set_verbosity(init_log_verbosity)

```

In []:

In []:

In []:

In []:

In []:

Notes Looking Forward to LoRA on RWKV

Hugging Face Community, seems to have a good portion of their models

<https://huggingface.co/RWKV>

<https://web.archive.org/web/20240530232509/https://huggingface.co/RWKV>

GitHub has even more versions/models, including the `v4-neo` that I think will be important (the LoRA project)

<https://github.com/BlinkDL/RWKV-LM/tree/main>

<https://web.archive.org/web/20240530232637/https://github.com/BlinkDL/RWKV-LM/tree/main>

The main RWKV website (?!)

<https://www.rwkv.com/>

<https://web.archive.org/web/20240529120904/https://www.rwkv.com/>

GOOD STUFF. A project doing LoRA with RWKV

<https://github.com/Blealtan/RWKV-LM-LoRA/>

<https://web.archive.org/web/20240530232823/https://github.com/Blealtan/RWKV-LM-LoRA>

The official blog, I guess, with some good coding examples

<https://huggingface.co/blog/rwkv>

<https://web.archive.org/web/20240530233025/https://huggingface.co/blog/rwkv>

It includes something that's similar to what I'm doing here in the `First_Full_LoRA_Trial_with_Transformer_Again.ipynb` tutorial, etc.

```
from transformers import AutoTokenizer, AutoModelForCausalLM

model_id = "RWKV/rwkv-raven-1b5"

model = AutoModelForCausalLM.from_pretrained(model_id).to(0)
tokenizer = AutoTokenizer.from_pretrained(model_id)
```

The `AutoModelForCausalLM` is the same as the tutorial I'm following, but I don't know what the `.to(0)` is for.

Really quickly, also looking at

<https://huggingface.co/RWKV/rwkv-4-world-7b>

<https://web.archive.org/web/20240530234438/https://huggingface.co/RWKV/rwkv-4-world-7b>

I see an example for CPU.

```
model = AutoModelForCausalLM.from_pretrained(
    "RWKV/rwkv-4-world-7b",
    trust_remote_code=True
).to(torch.float32)
```

```
tokenizer = AutoTokenizer.from_pretrained(  
    "RWKV/rwkv-4-world-7b",  
    trust_remote_code=True)
```

(Old version? Unofficial, it seems)

https://huggingface.co/docs/transformers/en/model_doc/rwkv

https://web.archive.org/web/20240530232341/https://huggingface.co/docs/transformers/en/model_doc/rwkv

In []: